

```

In [5]: import os, glob
import pandas as pd, numpy as np
import torch, torch.nn as nn, torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns

# --- DATA LOADER & PREPROCESSING ---
def load_and_merge_datasets(roots):
    merged = []
    for act_folder in sorted(os.listdir(roots[0])):
        if not act_folder.startswith('.'):
            sub_files = glob.glob(os.path.join(roots[0], act_folder, '*.csv'))
            for sub_file in sub_files:
                sub_name = os.path.basename(sub_file)
                triplet = []
                ok = True
                for r in roots:
                    path = os.path.join(r, act_folder, sub_name)
                    if os.path.exists(path):
                        triplet.append(pd.read_csv(path))
                    else:
                        ok = False
                        break
                if ok:
                    dfA, dfB, dfC = triplet
                    dfA = dfA.add_prefix('devicemotion_')
                    dfB = dfB.add_prefix('accelerometer_')
                    dfC = dfC.add_prefix('gyroscope_')
                    df = pd.concat([dfA, dfB, dfC], axis=1)
                    df['act'] = act_folder
                    merged.append(df)
            if not merged:
                raise FileNotFoundError("No data found. Check folder structure!")
    return pd.concat(merged, ignore_index=True)

def windows(df, window_size=100, stride=50):
    Xs, ys = [], []
    cols = [c for c in df.columns if c not in ['act']]
    arr = df[cols].values
    acts = df['act'].values
    for i in range(0, len(df) - window_size, stride):
        Xs.append(arr[i:i+window_size].T)
        ys.append(acts[i+window_size//2])
    return np.array(Xs), np.array(ys)

class MotionSense(Dataset):
    def __init__(self, X, y):
        self.X = torch.tensor(X, dtype=torch.float32)
        self.y = torch.tensor(y, dtype=torch.long)
    def __len__(self):
        return len(self.X)
    def __getitem__(self, i):
        return self.X[i], self.y[i]

class CNN_LSTM(nn.Module):
    def __init__(self, n_features, n_timesteps, n_classes):
        super().__init__()
        self.conv1 = nn.Conv1d(n_features, 64, 3, padding=1)

```

```

        self.bn1 = nn.BatchNorm1d(64)
        self.relu1 = nn.ReLU()
        self.conv2 = nn.Conv1d(64, 128, 3, padding=1)
        self.bn2 = nn.BatchNorm1d(128)
        self.relu2 = nn.ReLU()
        self.lstm = nn.LSTM(128, 64, batch_first=True, num_layers=2, dropout=0.3)
        self.fc1 = nn.Linear(64, n_classes)
        self.dropout = nn.Dropout(0.5)

    def forward(self, x):
        x = self.relu1(self.bn1(self.conv1(x)))
        x = self.relu2(self.bn2(self.conv2(x)))
        x = x.permute(0, 2, 1) # (batch, seq, feature)
        _, (h, _) = self.lstm(x)
        h = h[-1]
        h = self.dropout(h)
        return self.fc1(h)

def train_epoch(model, loader, criterion, optimizer, device):
    model.train()
    total_loss, correct, total = 0.0, 0, 0
    for Xb, yb in loader:
        Xb, yb = Xb.to(device), yb.to(device)
        optimizer.zero_grad()
        out = model(Xb)
        loss = criterion(out, yb)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
        optimizer.step()
        total_loss += loss.item() * Xb.size(0)
        pred = out.argmax(dim=1)
        correct += (pred == yb).sum().item()
        total += Xb.size(0)
    return total_loss / total, correct / total

def eval_loss(model, loader, criterion, device):
    model.eval()
    total_loss, correct, total = 0.0, 0, 0
    with torch.no_grad():
        for Xb, yb in loader:
            Xb, yb = Xb.to(device), yb.to(device)
            out = model(Xb)
            loss = criterion(out, yb)
            total_loss += loss.item() * Xb.size(0)
            pred = out.argmax(dim=1)
            correct += (pred == yb).sum().item()
            total += Xb.size(0)
    return total_loss / total, correct / total

def plot_history_pt(train_accuracies, test_accuracies, train_losses, test_losses, a
fig, axs = plt.subplots(1, 2, figsize=(16, 6))
axs[0].plot(train_accuracies, 'r', label='Train')
axs[0].plot(test_accuracies, 'b', label='Test')
axs[0].set_ylabel('Accuracy')
axs[0].set_xlabel('Epochs')
axs[0].set_title(add_title + f'Best Test Acc: {np.max(test_accuracies):.4f} @ {
axs[0].legend(); axs[0].grid()
axs[1].plot(train_losses, 'r', label='Train')
axs[1].plot(test_losses, 'b', label='Test')
axs[1].set_ylabel('Crossentropy Loss')
axs[1].set_xlabel('Epochs')
axs[1].legend(); axs[1].grid()
plt.tight_layout()
plt.show()

```

```

if __name__ == "__main__":
    print("Loading and merging data from all sensor folders...")
    roots = ['A_DeviceMotion_data', 'B_Accelerometer_data', 'C_Gyroscope_data']
    df = load_and_merge_datasets(roots)
    df = df.fillna(0)
    print(f"Loaded dataset shape: {df.shape}")

    print("\nSegmenting into windowed samples for deep learning...")
    X, y_str = windows(df, window_size=100, stride=50)
    print(f"Feature array shape: {X.shape}, Label array shape: {y_str.shape}")

    print("\nNormalizing and encoding the data...")
    nsamples, nfeat, nsteps = X.shape
    X_reshaped = X.transpose(0,2,1).reshape(-1, nfeat)
    X_reshaped = np.nan_to_num(X_reshaped, nan=0.0, posinf=0.0, neginf=0.0)
    scaler = StandardScaler()
    X_scaled = scaler.fit_transform(X_reshaped)
    X = X_scaled.reshape(nsamples, nsteps, nfeat).transpose(0,2,1)
    X = np.nan_to_num(X, nan=0.0, posinf=0.0, neginf=0.0)

    le = LabelEncoder()
    y_enc = le.fit_transform(y_str)
    print(f"Classes found: {list(le.classes_)}")

    X_tr, X_te, y_tr, y_te = train_test_split(X, y_enc, test_size=0.2, random_state=42)
    train_loader = DataLoader(MotionSense(X_tr, y_tr), batch_size=128, shuffle=True)
    test_loader = DataLoader(MotionSense(X_te, y_te), batch_size=128, shuffle=False)

    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    n_features = X_tr.shape[1]
    n_timesteps = X_tr.shape[2]
    n_classes = len(le.classes_)
    print("\nBuilding CNN-LSTM model...")
    model = CNN_LSTM(n_features, n_timesteps, n_classes).to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=1e-4)

    print("\nTraining model. For each epoch, loss and accuracy are reported for an epoch")
    epochs = 20
    train_losses, train_accuracies, test_accuracies, test_losses = [], [], [], []
    for ep in range(1, epochs+1):
        loss, train_acc = train_epoch(model, train_loader, criterion, optimizer, device)
        test_loss, test_acc = eval_loss(model, test_loader, criterion, device)
        train_losses.append(loss)
        train_accuracies.append(train_acc)
        test_accuracies.append(test_acc)
        test_losses.append(test_loss)
        print(f"Epoch {ep:02d} - Train Loss: {loss:.4f} - Train Acc: {train_acc:.4f}")

    print("\nPlotting training/validation accuracy and loss for each epoch. Use this plot to visualize the model's performance")
    plot_history_pt(train_accuracies, test_accuracies, train_losses, test_losses, epochs)

    print("\nEvaluating model on the test set for detailed classification performance")
    model.eval()
    all_labels, all_preds = [], []
    with torch.no_grad():
        for Xb, yb in test_loader:
            Xb = Xb.to(device)
            out = model(Xb)
            preds = out.argmax(dim=1).cpu().numpy()
            all_preds.extend(preds)
            all_labels.extend(yb.numpy())
    all_labels = np.array(all_labels)
    all_preds = np.array(all_preds)

```

```

print("\nClassification report below shows precision, recall, and F1-score for
print("These metrics help you understand per-activity performance, not just ove
print(classification_report(all_labels, all_preds, target_names=le.classes_))

print("The following confusion matrix heatmap visualizes where the model is mos
cm = confusion_matrix(all_labels, all_preds)
plt.figure(figsize=(12, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=le.classes_, yti
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.tight_layout()
plt.show()

print("Below are a few sample predictions from the test set (True Label vs Pred
print("This qualitative check lets you see actual predictions and errors.")
for i in range(10):
    true_label = le.classes_[all_labels[i]]
    pred_label = le.classes_[all_preds[i]]
    print(f"True: {true_label}, Pred: {pred_label}")

```

Loading and merging data from all sensor folders...  
Loaded dataset shape: (1433825, 22)

Segmenting into windowed samples for deep learning...  
Feature array shape: (28675, 21, 100), Label array shape: (28675,)


Normalizing and encoding the data...  
Classes found: ['dws\_1', 'dws\_11', 'dws\_2', 'jog\_16', 'jog\_9', 'sit\_13', 'sit\_5', 'std\_14', 'std\_6', 'ups\_12', 'ups\_3', 'ups\_4', 'wlk\_15', 'wlk\_7', 'wlk\_8']

Building CNN-LSTM model...

Training model. For each epoch, loss and accuracy are reported for analysis.

Epoch 01	– Train Loss: 2.3337	– Train Acc: 0.3349	– Test Loss: 1.7526	– Test Acc: 0.4600
Epoch 02	– Train Loss: 1.5619	– Train Acc: 0.4883	– Test Loss: 1.3019	– Test Acc: 0.5262
Epoch 03	– Train Loss: 1.2852	– Train Acc: 0.5476	– Test Loss: 1.1258	– Test Acc: 0.6115
Epoch 04	– Train Loss: 1.1561	– Train Acc: 0.5966	– Test Loss: 1.0093	– Test Acc: 0.6677
Epoch 05	– Train Loss: 1.0575	– Train Acc: 0.6434	– Test Loss: 0.9250	– Test Acc: 0.6980
Epoch 06	– Train Loss: 0.9871	– Train Acc: 0.6716	– Test Loss: 0.8631	– Test Acc: 0.7247
Epoch 07	– Train Loss: 0.9154	– Train Acc: 0.6997	– Test Loss: 0.8007	– Test Acc: 0.7388
Epoch 08	– Train Loss: 0.8585	– Train Acc: 0.7177	– Test Loss: 0.7278	– Test Acc: 0.7658
Epoch 09	– Train Loss: 0.8248	– Train Acc: 0.7322	– Test Loss: 0.6812	– Test Acc: 0.7801
Epoch 10	– Train Loss: 0.7861	– Train Acc: 0.7442	– Test Loss: 0.6595	– Test Acc: 0.7833
Epoch 11	– Train Loss: 0.7497	– Train Acc: 0.7533	– Test Loss: 0.6312	– Test Acc: 0.7888
Epoch 12	– Train Loss: 0.7194	– Train Acc: 0.7586	– Test Loss: 0.6161	– Test Acc: 0.7934
Epoch 13	– Train Loss: 0.7000	– Train Acc: 0.7665	– Test Loss: 0.5727	– Test Acc: 0.8023
Epoch 14	– Train Loss: 0.6705	– Train Acc: 0.7743	– Test Loss: 0.5552	– Test Acc: 0.8000
Epoch 15	– Train Loss: 0.6559	– Train Acc: 0.7747	– Test Loss: 0.5322	– Test Acc: 0.8117
Epoch 16	– Train Loss: 0.6353	– Train Acc: 0.7798	– Test Loss: 0.5189	– Test Acc: 0.8120
Epoch 17	– Train Loss: 0.6116	– Train Acc: 0.7856	– Test Loss: 0.5208	– Test Acc: 0.8087
Epoch 18	– Train Loss: 0.6118	– Train Acc: 0.7818	– Test Loss: 0.5154	– Test Acc: 0.8061
Epoch 19	– Train Loss: 0.5904	– Train Acc: 0.7877	– Test Loss: 0.4927	– Test Acc: 0.8176
Epoch 20	– Train Loss: 0.5755	– Train Acc: 0.7977	– Test Loss: 0.4734	– Test Acc: 0.8185

Plotting training/validation accuracy and loss for each epoch. Use this to visually assess overfitting or convergence.




Evaluating model on the test set for detailed classification performance...

Classification report below shows precision, recall, and F1-score for each class. These metrics help you understand per-activity performance, not just overall accuracy.

	precision	recall	f1-score	support
dws_1	0.41	0.49	0.45	204
dws_11	0.00	0.00	0.00	92
dws_2	0.53	0.64	0.58	241
jog_16	0.66	0.17	0.27	122
jog_9	0.81	0.96	0.88	423
sit_13	0.90	0.93	0.92	422
sit_5	0.97	0.95	0.96	951
std_14	0.82	0.84	0.83	331
std_6	0.94	0.93	0.94	911
ups_12	0.00	0.00	0.00	115
ups_3	0.64	0.03	0.06	237
ups_4	0.44	0.95	0.60	289
wlk_15	0.96	0.95	0.95	266
wlk_7	0.98	0.89	0.93	643
wlk_8	0.85	0.98	0.91	488
accuracy			0.82	5735
macro avg	0.66	0.65	0.62	5735
weighted avg	0.81	0.82	0.79	5735

The following confusion matrix heatmap visualizes where the model is most accurate and which classes are being confused with others.

```
C:\Users\Sachin\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:134
4: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to control
this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Sachin\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:134
4: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to control
this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Sachin\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:134
4: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to control
this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```



Below are a few sample predictions from the test set (True Label vs Predicted Label).

This qualitative check lets you see actual predictions and errors.

True: ups\_4, Pred: ups\_4  
 True: wlk\_8, Pred: wlk\_8  
 True: jog\_16, Pred: jog\_9  
 True: sit\_5, Pred: sit\_5  
 True: std\_6, Pred: std\_6  
 True: std\_6, Pred: std\_6  
 True: wlk\_15, Pred: wlk\_15  
 True: ups\_3, Pred: ups\_4  
 True: sit\_5, Pred: sit\_5  
 True: wlk\_8, Pred: wlk\_8

In [ ]:

In [ ]: